Dataset Analysis for Decision-Making

27/03/2025

HIV_Epidemiology.... .csv
19.2 MB

Analyze the provided dataset to identify its key characteristics for data-driven decision-making. Upon receiving a specific analytical question, address it strictly using the dataset. Your deliverables must include a clear, concise answer to the question and the Python code that extracts and processes the data used to derive your answer.

qwen-max-latest

To analyze the dataset and address your request, I will first load the dataset (`HIV_Epidemiology_Children_Adolescents_2024_out.csv`) into Python using `pandas` . Then, I will explore its structure, identify key characteristics, and prepare to answer any specific analytical questions you may have.

Step 1: Load and Inspect the Dataset

I'll start by loading the dataset and inspecting its structure (e.g., column names, data types, missing values, and summary statistics). This will help us understand the dataset's key characteristics.

python

```python
import pandas as pd

# Load the dataset
file_path = 'HIV_Epidemiology_Children_Adolescents_2024_out.csv'
data = pd.read_csv(file_path)

# Display basic information about the dataset
print("Dataset Overview:")
print(data.info())

# Display the first few rows of the dataset
print("\nFirst 5 Rows of the Dataset:")
print(data.head())

# Check for missing values
print("\nMissing Values per Column:")
print(data.isnull().sum())

# Summary statistics for numerical columns
print("\nSummary Statistics:")
print(data.describe())
```

Step 2: Identify Key Characteristics

From the output of the above code, we can identify:

1. **Columns** : The dataset contains variables such as demographics (e.g., age, gender), epidemiological data (e.g., HIV prevalence, incidence), and geographic or temporal information.
2. **Data Types** : Numerical, categorical, and possibly datetime columns.
3. **Missing Data** : Any missing values that need handling.
4. **Summary Statistics** : Mean, median, min, max, and other descriptive statistics for numerical variables.

Step 3: Address Specific Analytical Questions

Once you provide a specific analytical question, I will write Python code to extract and process the relevant data from the dataset to answer it. For example:

- **Question Example** : "What is the average HIV prevalence among children aged 0–14 in the dataset?"
- **Answer** : I would filter the dataset for the relevant age group, calculate the mean HIV prevalence, and provide the result along with the corresponding Python code.

Please provide your specific analytical question so I can proceed with the analysis.

Propose three appropriate statistical tests and select the best one to determine whether the observed differences in HIV incidence rates between adolescent girls and boys in Europe over time differ from those in Africa. Then, apply the selected test and extract the results.

qwen-max-latest

To address the question of whether the observed differences in HIV incidence rates between adolescent girls and boys in Europe over time differ from those in Africa, we need to carefully consider the structure of the dataset and choose appropriate statistical tests. Below is a detailed explanation of the proposed tests, the rationale for selecting the best one, and the implementation of the selected test.

Step 1: Understanding the Dataset

From the dataset (`HIV_Epidemiology_Children_Adolescents_2024_out.csv`), we expect columns such as:

- `Region` : Indicates geographic regions (e.g., Europe, Africa).
- `Gender` : Specifies the gender of adolescents (e.g., "Girls", "Boys").
- `Year` : Temporal information.
- `HIV_Incidence_Rate` : Numerical values representing HIV incidence rates.

We will:

1. Filter the data for "Europe" and "Africa".
2. Compare the differences in HIV incidence rates between adolescent girls and boys across these two regions over

time.

Step 2: Proposed Statistical Tests

Here are three appropriate statistical tests for this analysis:

1. Two-Way ANOVA

- **Purpose** : To determine if there is a significant interaction effect between region (Europe vs. Africa) and gender (girls vs. boys) on HIV incidence rates over time.
- **Assumptions** :
  - Normality of residuals.
  - Homogeneity of variances.
  - Independence of observations.
- **Why it fits** : This test allows us to examine the interaction between two categorical variables (region and gender) on a continuous outcome (HIV incidence rate).

2. Mixed-Effects Model

- **Purpose** : To account for repeated measures over time (if multiple years are available) while testing the interaction between region and gender.
- **Advantages** :
  - Handles hierarchical or longitudinal data.
  - Can include random effects (e.g., subject-specific variability).
- **Why it fits** : If the dataset contains repeated observations over time for the same individuals or groups, this model is more robust.

3. Mann-Whitney U Test (Non-Parametric)

- **Purpose** : To compare the distribution of HIV incidence rates between two independent groups (e.g., Europe vs. Africa) without assuming normality.
- **Assumptions** :
  - Data does not need to follow a normal distribution.
- **Why it fits** : If the data violates assumptions of normality or homogeneity of variances, this test can be used as a fallback.

Step 3: Selecting the Best Test

The **Two-Way ANOVA** is the most appropriate test here because:

1. It directly tests the interaction between region and gender, which is the primary research question.
2. It assumes normality and homogeneity of variances, which can be checked using diagnostic plots.
3. It provides clear insights into both main effects (region and gender) and their interaction.

If the assumptions of ANOVA are violated, we can fall back to the Mann-Whitney U Test.

Step 4: Implementing the Two-Way ANOVA

Below is the Python code to perform the analysis:

python
```python
import pandas as pd
import statsmodels.api as sm
from statsmodels.formula.api import ols
import matplotlib.pyplot as plt
import seaborn as sns

# Load the dataset
file_path = 'HIV_Epidemiology_Children_Adolescents_2024_out.csv'
data = pd.read_csv(file_path)

# Filter data for Europe and Africa, and adolescent age group
filtered_data = data[(data['Region'].isin(['Europe', 'Africa'])) &
                     (data['Age_Group'] == 'Adolescents')]

# Check for missing values
filtered_data.dropna(subset=['HIV_Incidence_Rate', 'Region', 'Gender'], inplace=True)

# Perform Two-Way ANOVA
model = ols('HIV_Incidence_Rate ~ C(Region) * C(Gender)', data=filtered_data).fit()
anova_table = sm.stats.anova_lm(model, typ=2)

# Display the ANOVA results
print("Two-Way ANOVA Results:")
print(anova_table)

# Visualize the interaction effect
sns.pointplot(data=filtered_data, x='Region', y='HIV_Incidence_Rate', hue='Gender', err
plt.title('Interaction Effect: Region x Gender on HIV Incidence Rates')
plt.ylabel('HIV Incidence Rate')
plt.xlabel('Region')
plt.show()
```

Step 5: Interpretation of Results

1. **ANOVA Table** :
   - The `C(Region)` row tests the main effect of region.
   - The `C(Gender)` row tests the main effect of gender.
   - The `C(Region):C(Gender)` row tests the interaction effect between region and gender.
   - A significant p-value (< 0.05) for the interaction term indicates that the difference in HIV incidence rates between girls and boys varies significantly between Europe and Africa.
2. **Visualization** :
   - The point plot shows the mean HIV incidence rates for girls and boys in Europe and Africa, with error bars indicating standard errors.

- If the lines for girls and boys cross or diverge significantly between regions, it suggests an interaction effect.

Final Answer

After running the above code, the output of the ANOVA table will provide the p-values for the interaction term. If the p-value for `C(Region):C(Gender)` is significant (e.g., < 0.05), we conclude that the observed differences in HIV incidence rates between adolescent girls and boys differ significantly between Europe and Africa.

Please run the code and share the results if you'd like further interpretation!